

Dylan Encarnacion Pell

Mechanical Engineering · Liquid Propulsion & Fluid Systems

I design, build, and test liquid rocket propulsion hardware. My work has moved from single components, to owning the feed system on a vehicle that flew, to leading the architecture of a full cryogenic fluid system.

Los Angeles, CA d.pell016@gmail.com

This portfolio is organized as a timeline. Each project records the problem, the constraints, the reasoning behind the design decisions, and how the result was validated. The through line is a progression in liquid propulsion, from component design, to owning the feed system on a vehicle that flew, to leading full-system architecture on a cryogenic stand. Alongside the hardware are the analysis tools I built for my teams to design against, and a professional R&D internship that applied the same discipline, defining what good enough means and proving it, to a sourcing decision rather than a rocket.

FS-03	Project Clementine	2026 - Present
TM-01	Project Turbopump	2026 - Present
FS-02	Project Poseidon	2025 - 2026
SW-01	Fluid System Sizing & Trade Study Tool	2025 - Present
RD-01	Supplier Cost-Benefit Analysis & Validation	Jul 2025 - Oct 2025
MF-01	Bulkhead Drill Jig	2025
ST-01	Composite Overwrapped Pressure Vessel Study	2024 - 2025
MD-01	Formula SAE Suspension Upright	2024 - 2025

Project Clementine

LOX / Ethanol Liquid Bipropellant Rocket

ROLE	ORGANIZATION	PERIOD
Vice President; previously Lead Fluid Systems Engineer	Highlander Space Program, UC Riverside	2026 - Present




Figure 1. The LOX and ethanol fluid wall, partway through assembly.

The problem

Our previous vehicle, [Project Poseidon](#), used nitrous oxide, which is self-pressurizing: the propellant supplies its own tank pressure through its vapor pressure, so there is no separate pressurization system to design. Moving to liquid oxygen removes that convenience. LOX has to be pushed, which means the whole pressurization architecture becomes something I had to design rather than inherit.

I led the P&ID and did the bulk of the design myself, with the wider fluids team contributing as the work came together. I sized the lines, calculated the pressure drops, and made sure the system delivered pressure and flow where they needed to go. I selected the pneumatics, valves, and regulators, set the pressure schedule, defined the interfaces between the pressurization, propellant, and ignition subsystems, presented the updates to the team every week, and installed the result on our fluids test wall. The P&ID itself belongs to Highlander Space Program.

Constraints

- **Single pressurant source.** One GN2 bottle had to serve both the pneumatic actuation and the propellant tanks. Budget did not allow separate supplies.
 - **Cryogenic propellant.** LOX introduces material compatibility, thermal, and trapped volume hazards that a room-temperature system does not have.
 - **Student build.** Commercial off-the-shelf components, limited machining, and a team where members rotate in and need to operate the system safely without deep context.
 - **Every test is a hazardous operation.** The system had to be safe to operate and legible to a new operator, not just correct on paper.
-

Design decisions

Holding tank pressure constant as the bottle drains

A single GN2 bottle blows down as gas is consumed. Treating the bottle as ideal gas,

$$\frac{p_1 V}{T_1} = \frac{p_2 V}{T_2}$$

supply pressure falls continuously through a test. A single spring regulator tracks that decay: spring regulators exhibit **droop**, where outlet pressure falls as inlet pressure drops and as flow increases, because the spring force is fixed while the sensing area sees changing conditions.

Inconsistent tank pressure means inconsistent chamber pressure, which means the engine is not being tested at the condition it was designed for.

My solution was a **two-stage regulation scheme**: a spring regulator sets the reference pressure in the dome of a dome-loaded regulator. The dome regulator references that near-constant dome pressure across a large diaphragm area rather than a spring, so its outlet stays essentially flat as the bottle decays and as flow varies. The spring regulator only has to hold a low-flow reference, which is the job it does well.

Sizing the pressurant lines

During checkout I measured a pressure drop across the GN2 lines that was larger than the budget allowed. Rather than raise the supply pressure to compensate, I traced it to undersized tubing and corrected the geometry.

The reasoning is the Darcy-Weisbach relation for major losses,

$$\Delta p = f \frac{L}{D} \frac{\rho v^2}{2}$$

with the flow velocity set by continuity,

$$v = \frac{4Q}{\pi D^2}$$

Substituting velocity into the pressure drop shows why line diameter dominates:

$$\Delta p \propto \frac{1}{D} \cdot \frac{1}{D^4} = \frac{1}{D^5}$$

At fixed volumetric flow, pressure drop scales with the **inverse fifth power** of internal diameter. Going up a single tube size is worth far more than any amount of tuning elsewhere in the run, and it fixes the problem at its source instead of masking it with higher supply pressure. Upsizing the GN2 lines brought the drop back inside budget.

Fittings, valves, and bends were then added as minor losses,

$$\Delta p_{\text{minor}} = K \frac{\rho v^2}{2}, \quad \Delta p_{\text{total}} = \sum \Delta p_{\text{major}} + \sum \Delta p_{\text{minor}}$$

with the friction factor from the Reynolds number,

$$Re = \frac{\rho v D}{\mu}$$

Relief valves on every isolable section

Any volume that can be closed off at both ends by ball valves can trap fluid. With a cryogenic system this is not a theoretical concern: trapped liquid oxygen warming toward ambient expands by roughly a factor of 860 from liquid to gas, and in a rigid closed volume that becomes a pressure rise that will find the weakest component.

I walked the P&ID section by section, identified every volume that could be isolated between two valves, and placed a relief valve on each one. The rule I designed to was that no operator sequence, correct or incorrect, should be able to create a sealed volume with no path to relieve.

Instrumentation at every node

Each measurement point carries **both** a mechanical pressure gauge and a pressure transducer. This is deliberate redundancy with two different failure modes:

- The **transducer** feeds the DAQ, giving a continuous recorded trace for post-test analysis and live monitoring.
- The **gauge** is independent of power and software. An operator can confirm the state of the system by eye during a walkdown, and a disagreement between gauge and transducer immediately separates an instrumentation fault from a real fluid anomaly.

That distinction matters during a no-go: knowing whether the sensor is lying or the system is genuinely off-nominal determines whether you safe the stand or keep working.

Igniter propellant split

The tanks divert a small portion of fuel and oxidizer to a torch igniter that lights the main combustion chamber. Sizing that bleed is an orifice problem,

$$\dot{m} = C_d A \sqrt{2\rho \Delta p}$$

The split has to deliver a reliable igniter flow without meaningfully disturbing the pressure or flow delivered to the main injector.

What I built

A complete pressurization, propellant, and ignition fluid system: GN2 supply feeding both pneumatic actuation and tank pressurization through two-stage regulation, run tanks for fuel and oxidizer, poppet pneumatic valve actuation, a torch igniter feed, relief protection on every isolable section, and instrumentation at each node reporting to the DAQ. The system is installed on our fluids test wall and is the stand the team now tests on.

Validation

1. **Component bench testing** before integration, to confirm each regulator, valve, and transducer behaved as specified rather than as assumed.
2. **System pressure checks** using the installed transducers, comparing measured drops against the values predicted by the sizing model.
3. **Correction and re-test:** the GN2 line pressure drop was identified here, corrected by upsizing, and confirmed against the model.

What I would do differently

I would build the pressure drop budget before selecting tubing rather than confirming it during checkout. The physics was not a surprise; I had the sizing tool that computes it. Running the numbers first would have caught the undersized run on paper and saved a rework cycle on hardware. It is the clearest lesson I have taken from this system: analysis is cheapest before you buy the fittings.




Figure 2. An early draft of the P&ID, iteration 5.




Figure 3. Iteration 6.




Figure 4. Iteration 7, as the architecture converged.



Figure 5. Drafting the fluid wall layout on a plank of wood, before the real build.

Project Turbopump

Electric-drive propellant pumps, from first principles to first prints

ROLE	ORGANIZATION	PERIOD
Sole designer	Personal project	2026 - Present

The problem

Project Clementine is pressure-fed: high-pressure GN2 pushes on the LOX and ethanol to move them through the system. It works, but the gas and the vessels that hold it are dead weight, and tank pressure caps how hard I can run the engine. I wanted to know if I could pump instead. Two small motor-driven pumps, one per propellant, that add the pressure actively instead of storing it. That is the electric pump-fed idea, the one Rocket Lab flies on Rutherford, and I wanted to see if I could design it myself. So I started a personal project, branched it off Clementine, and worked backward from the engine: an 8.8 kN ethanol/LOX engine running 600 psig chamber pressure.

One point of precision, because a good reviewer will check it: a "turbopump" strictly means a pump spun by a turbine. What I designed is the pump end, which is identical either way, driven by an electric motor for now. Turbine drive is a planned later phase. I keep the name Project Turbopump because that is where it is headed.

Constraints

- **First principles, alone.** No team, no inherited design. If I could not derive a number and cite where the method came from, I did not use it.
- **Cryogenic oxygen service.** LOX is not water. Material compatibility, cavitation, and cleanliness all come with it, held to oxygen-service standards (CGA G-4.1, ASTM G93).

Deriving the requirements

Everything starts at the engine. From thrust and specific impulse I worked back to mass flow,

$$\dot{m} = \frac{F}{I_{sp} g_0}$$

and from there to what each pump has to deliver: 441 m of head on the LOX side and 806 m on the ethanol, which also has to feed the regen jacket, at 1.86 and 1.68 L/s, for 30 to 44 kW of shaft power. The first useful result was not a calculation, it was a catch. My first-principles mass flow disagreed with the

data sheet I was handed, and the reason was a units-label error on the sheet: it read lbm/s for a number that was really kg/s. That set the rule for the whole project. Derive it myself, then trust it.

The architecture, and why

Three calls shaped everything downstream, and I made each to defend, not to assume.

Electric drive, not a gas-generator turbine. A real turbopump burns propellant to spin a turbine. For one person to build and control, a motor is far simpler, and it flies now, so I get the simplicity without the hand-waving. The pump end carries straight over to a turbine later, so nothing here is wasted.

Two independent pumps, one per propellant, instead of one shaft driving both. LOX and ethanol want different speeds and behave differently, and splitting them kills a whole class of compromise.

No inducer. My NPSH analysis was the make-or-break of the project. LOX comes in at 90.2 K, which is its boiling point at one atmosphere, so at ambient tank pressure the suction margin is actually negative and the pump would flash to vapor before it did any work. The textbook fix is an inducer, a little screw ahead of the impeller, and it is also the hardest part to machine. Instead of adding it, I turned the finding into a hard requirement, pad the LOX tank to 3 to 4 bar, and deleted the riskiest part of the pump. I spent tank pressure to buy the margin back.

The LOX pump

The first-cut LOX impeller is a 20,000 rpm, 83.6 mm semi-open centrifugal wheel: six blades, 25 degrees of backsweep, Wiesner slip, and an inlet eye sized to hold off cavitation, closing self-consistently to the 441 m the engine needs. I went semi-open because a closed impeller with 1.4 mm internal passages is neither machinable nor inspectable at student level, and inspectable matters when the oxygen-cleanliness rules make me verify every wetted surface.

The number that scared me was axial thrust, near-discharge pressure on the back of the impeller trying to shove the rotor into its housing. My first pass put it around 8 kN. A later self-audit against my sizing tool caught that I had integrated the pressure over the wrong radius, and the real figure is closer to 15 kN. Either way the fix holds: a back wear ring plus six balance holes bleed the back face down and drop the residual to a few hundred newtons, which an angular-contact bearing pair carries easily. I sized the shaft for a first bending critical near 66,000 rpm, a 3.3x margin over running speed.

The whole pump is parametric CAD in Python (CadQuery), so every part regenerates from the inputs and exports to STEP and STL. Change one number and the pump redraws. I also built a pump sizing tool in a spreadsheet, a few hundred live formulas, that runs the entire chain from engine data to impeller geometry, cavitation margin, and the minimum tank pressure, and flags when a design lands in the un-machinable corner. It reproduces this pump exactly, and I can point it at any pump I design next.

Problems I caught and fixed

This is the part I most want to show, because it is where the engineering actually happened.

- **The ethanol pump passed every analysis, then failed at the drawing board.** On paper it closed. In geometry, its very low specific speed drove the impeller exit width to 0.7 mm, too narrow to machine or to flow cleanly. I did not force it. I laid out three escape routes, spin it faster, split it into two stages, or switch to a partial-emission design, and gated the choice on which motor I can actually buy.
- **I found a cheap motor that looked perfect and rejected it.** A hundred-dollar NEMA 17, until a power-balance check, $P = \tau \omega$, showed it was about 200 times undersized. The useful part was the one-line rule I pulled out of it, so now I can size any motor before spending a cent. The real target is 8 to 9 N·m continuous, an eVTOL-class inrunner.
- **I keep catching my own work with a second method.** Re-optimizing the inlet eye invalidated a blade wrap angle I had already called done, so I redid it. A self-review of the first CAD caught a volute sized 2.6 times too big for the flow. And the thrust-radius error above came out of auditing the design doc against my sizing tool. None of them shipped, because I check every number twice.
- **The first 3D prints caught three bugs the math never would have.** When I printed the parts and dry-fit them, I found a discharge bore that stopped just short of breaking through the casing wall and pointed against the swirl, a bearing housing I had never actually modeled, and the worst one: the impeller perched on its own blade tips because I had cut a flat pocket floor against blades that taper from 8 mm tall at the eye to 1.4 mm at the rim. In a semi-open pump the stationary wall has to be contoured to the blade profile; it is part of the flow passage, not a lid. All three are fixed in the current CAD. Holding the part in your hand finds what the screen hides.

Validation

I cannot hot-fire a LOX pump in a garage, so I wrote a water-similarity test program to prove the design safely. It runs the pump on water and scales the results with the affinity laws, collapsing head and flow onto their dimensionless coefficients so that water data predicts the LOX behavior. Staged the right way it needs only about 4% of full power, the same route ESA and published LOX/methane electric-pump programs use to de-risk on the bench. I wrote the procedure with abort criteria and containment from the start, and priced the whole rig at roughly 600 to 1,100 dollars in parts.

Status and what is next

I am keeping this honest about where it stands. The LOX pump design, the sizing tool, and the test program are done, and I have the first 3D-printed parts in hand and dry-fitting, which is where the last three bugs came from. Nothing is in metal yet, and the ethanol pump has one open decision waiting on which motor I source. Next is sourcing the drive, correcting the prints, machining the first aluminum, and running the water test. Every number traces to a derivation and every decision to a logbook entry, so this is built to be picked up and pushed forward, not admired as finished.

Project Poseidon

N₂O / Ethanol Liquid Bipropellant Rocket

ROLE	ORGANIZATION	PERIOD
Lead Fluid Systems Engineer	Highlander Space Program, UC Riverside	2025 - 2026



Figure 1. Poseidon on the launch rail before flight.

COMPETITION RESULT

1st place, Most Efficient Liquid Engine

EVENT

2025 FAR-OUT Advanced Propulsion Rocketry Competition

FLIGHT ALTITUDE

7,340 ft

The problem

Poseidon was UC Riverside's first liquid bipropellant rocket. Nothing existed to inherit: no prior feed system, no validated pressure vessel, no established test procedure. As Lead Fluid Systems Engineer I built the pressure vessel, and I designed, built, and validated the integrated fuel and oxidizer feed system.

Constraints

- **First of its kind for the program.** No internal heritage hardware or prior data to design against.

- **Competition deadline.** A fixed date that testing had to be complete before.
- **Student budget and shop access.** Design had to be manufacturable with the machining and assembly capability actually available to us.
- **Safety margin on a pressurized structure** with people standing near it.

Design decisions

Why nitrous oxide simplified the pressurization problem

N₂O is self-pressurizing. At room temperature its vapor pressure sits near 700 to 750 psi, so the propellant supplies its own tank pressure without a separate pressurant system. For a first liquid vehicle this removed an entire subsystem and let the team concentrate on the engine, the structure, and the feed path.

The tradeoff is that vapor pressure is strongly temperature dependent, so tank pressure, and therefore engine performance, varies with ambient conditions on test day. That sensitivity is the reason our next vehicle moved to a separately pressurized architecture, which became [Project Clementine](#).

Sizing the pressure vessel

I did not originate the vessel's structural analysis, but I double-checked it before we committed to building the hardware. The vessel was analyzed as a thin-walled cylinder, valid where the radius to thickness ratio is greater than about 10. Under internal pressure the wall carries a hoop stress and an axial stress,

$$\sigma_h = \frac{pr}{t}, \quad \sigma_a = \frac{pr}{2t}$$

Hoop stress is twice axial stress, so the hoop direction governs and sets the required wall thickness,

$$t \geq \frac{p_{\text{design}} r}{\sigma_{\text{allow}}}, \quad \sigma_{\text{allow}} = \frac{\sigma_y}{FS}$$

Design pressure was set from the maximum expected operating pressure with margin applied above it, rather than from nominal operating pressure, so that the structure is sized for the worst credible case rather than the intended one. The combined stress state was checked against yield using the von Mises criterion,

$$\sigma_v = \sqrt{\sigma_h^2 - \sigma_h \sigma_a + \sigma_a^2}$$

Material selection traded strength to weight against weldability and oxidizer compatibility, and the geometry was iterated in CAD against what the shop could actually produce and inspect.

Designing the feed path

The feed system had to deliver both propellants to the engine at the correct flow and pressure, sequence safely through fill and fire, and seal reliably. Line and orifice sizing followed the same incompressible relations used across the program,

$$\dot{m} = C_d A \sqrt{2\rho \Delta p}$$

with major and minor losses accumulated along each run to confirm the engine saw its design inlet condition rather than whatever the plumbing happened to deliver.

What I built

The pressure vessel, which I built and fabricated, coordinating with manufacturing until it was real hardware. Then the integrated fuel and oxidizer feed system, assembling and validating the feed lines, valves, and interfaces that connected tanks to engine.

Validation

Testing followed a deliberate ladder, each step qualifying the next:

1. **Leak and pressure checks** on the assembled system, to prove sealing before introducing propellant.
2. **Cold flow testing**, flowing propellant through the complete feed path without combustion. This validates flow rates, sequencing, and valve behavior while the consequences of being wrong are still small.
3. **Static fire**, running the full engine with the vehicle restrained, to confirm the feed system delivered propellant correctly under real operating and thermal conditions.
4. **Flight**.

The vehicle earned first place for Most Efficient Liquid Engine at the 2025 FAR-OUT Advanced Propulsion Rocketry Competition and reached 7,340 feet.

The pressure vessel that flew was drilled the night before the hot fire using a jig I designed and manufactured on the spot. That story is written up separately in [Bulkhead Drill Jig](#).

What I would do differently

I would instrument more heavily. We proved the system worked, but with more transducers across the feed path we would have had a far richer picture of *how* it worked, and a dataset to design the next vehicle against instead of starting Clementine's numbers from first principles again. The instrumentation density on Clementine is a direct response to this.



Figure 2. The pressure vessel I built and fabricated.



Figure 3. Integrated fuel and oxidizer feed system.



Figure 4. Static fire test campaign.



Figure 5. Launch. The vehicle reached 7,340 ft.

Fluid System Sizing & Trade Study Tool

Python / MATLAB design automation

ROLE	ORGANIZATION	PERIOD
Author	Highlander Space Program, UC Riverside	2025 - Present

LOX / IPA(or Ethanol) Engine — Auto Sizing									
Inputs (Edit the input blue cells)									
Parameter	Value	Units	Parameter	Value	Units	Parameter	Value	Units	Notes
Thrust (Sea Level)	2000	lbf	Run time per day	10	s	Number of days	1.5	-	Typically 2
WSP	1000	psi	WSP	100	psi	WSP	0.005	lbf	Applied to both legs
LOX density ρ_{LOX}	1143	kg/m ³	IPA density ρ_{IPA}	786	kg/m ³	Total mass	1.4438	kg	Mass engine, line, tank
LOX mass flow $\dot{m}_{LOX}$	2.527	kg/s	IPA mass flow $\dot{m}_{IPA}$	1.2403	kg/s	Line length L	2	m	Straight run minimum
LO fittings	0	-	Curve (Based on Reynolds number and schedule number)	0.05	-	Page ID	0.0158	m	See note 1.0.01005 m
Engine SP per valve	5	psi	Valve Cv (Based on schedule number and schedule number)	35	-	Hard ramp time	20	s	See note 1.0.01005 m
Outputs									
Quantity	LOX	Ethanol	Units	Formula (left)	Equation (right)				
Mass per line day	21.27	13.203	kg	$\dot{m}_1 = \dot{m}_2$	$\dot{m}_1 = \dot{m}_2$				
Mass line to flow rate	31.305	19.305	kg	$\dot{m}_1 = \dot{m}_2$	$\dot{m}_1 = \dot{m}_2$				
Liquid volume V_{liq}	0.0281022253	0.024451086	m ³	$V_{liq} = \frac{\dot{m}_{liq}}{\rho_{liq}} (1 + \beta)$	$V_{liq} = \frac{\dot{m}_{liq}}{\rho_{liq}} (1 + \beta)$				
Tank volume V_{tank}	0.032244887	0.0282788296	m ³	$V_{tank} = \frac{V_{liq}}{1 - \beta}$	$V_{tank} = \frac{V_{liq}}{1 - \beta}$				
Tank volume (L)	31.2248817	28.2788208	L	$V_{tank} = 1000 V_{liq}$	$V_{tank} = 1000 V_{liq}$				
Tank volume (gal US)	8.18888247	7.463972101	gal	$V_{tank} = \frac{V_{liq}}{0.001308}$	$V_{tank} = \frac{V_{liq}}{0.001308}$				
Volumetric flow Q	0.00188415425	0.00188818735	m ³ /s	$Q = \dot{m} / \rho$	$Q = \dot{m} / \rho$				
Volumetric flow (gpm)	20.54744059	26.76887011	gpm	$Q_{gpm} = Q \cdot 15850.323$	$Q_{gpm} = Q \cdot 15850.323$				
Velocity V	5.507750759	6.610267613	m/s	$V = \frac{Q}{A}$	$V = \frac{Q}{A}$				
Dynamic pressure p_{dyn}	51.71.87406	23.15.34787	Pa	$p_{dyn} = \frac{\rho V^2}{2}$	$p_{dyn} = \frac{\rho V^2}{2}$				
SP (gpm) (curve)	0.0403.0004	0.0403.0004	Pa	$\Delta p_{dyn} = f \cdot \frac{\rho V^2}{2}$	$\Delta p_{dyn} = f \cdot \frac{\rho V^2}{2}$				
SP (fittings) (K)	0	0	Pa	$\Delta p_K = \sum K \cdot \frac{\rho V^2}{2}$	$\Delta p_K = \sum K \cdot \frac{\rho V^2}{2}$				
Total SP (Total Pressure Drop across the line)	26.74481318	26.74481318	psi	$\Delta p_{tot} = (\Delta p_{dyn} + \Delta p_K) \cdot \rho / \rho_{liq}$	$\Delta p_{tot} = (\Delta p_{dyn} + \Delta p_K) \cdot \rho / \rho_{liq}$				
Required C_v	14.11480124	10.6313846	-	$C_v = Q_{gpm} \sqrt{\frac{SG}{\Delta p_{valve}}}$	$C_v = Q_{gpm} \sqrt{\frac{SG}{\Delta p_{valve}}}$				
Valve SP with C_v (curve)	0.8131812268	0.4150101033	psi	$\Delta p_{valve} = \left(\frac{C_v}{Q_{gpm}}\right)^2 SG$	$\Delta p_{valve} = \left(\frac{C_v}{Q_{gpm}}\right)^2 SG$				
Summary (Not fully automated)									
Item	LOX	IPA	Units	Comment					
Run tank minimum volume	31.2248817	28.2788208	L	1000 L tank					
Line size (diameter)	3/2, 3/4, 1/2	1/2, 3/4, 1/2	-	Line size (diameter) (diameter)					
Line SP (Pressure Drop)	26.74481318	26.74481318	psi	Assume line is 1000 L tank					
Required valve C_v	14.11480124	10.6313846	-	Required valve C_v					
SP with 1" full port C_v	0.8131812268	0.4150101033	psi	Typical $C_v = 10$ (diameter)					

Figure 1. Tool output for tank and line sizing with pressure-drop calculations.

The problem

Every change to a feed system propagates. Change a tube size and the velocity, Reynolds number, friction factor, and pressure drop all move, which changes the tank pressure you need, which changes the regulator setpoint and the structural requirement on the tank. Doing that by hand takes long enough that in practice you stop exploring alternatives and commit to the first workable answer.

I wanted the team to be able to ask "what if this line were larger" and get an answer immediately, so that design decisions came from comparison rather than from whichever configuration someone happened to calculate first.

What it computes

Given propellant properties, target mass flow, and a proposed line and tank geometry, the tool returns total pressure drop, flow rates, and line and tank sizing.

Flow velocity from continuity:

$$v = \frac{4Q}{\pi D^2}$$

Reynolds number to establish the flow regime:

$$Re = \frac{\rho v D}{\mu}$$

Friction factor from the Swamee-Jain explicit approximation, which avoids iterating the implicit Colebrook equation:

$$f = \frac{0.25}{\left[\log_{10} \left(\frac{\varepsilon}{3.7D} + \frac{5.74}{Re^{0.9}} \right) \right]^2}$$

Major losses along each run:

$$\Delta p_{\text{major}} = f \frac{L}{D} \frac{\rho v^2}{2}$$

Minor losses through valves, elbows, and fittings:

$$\Delta p_{\text{minor}} = K \frac{\rho v^2}{2}$$

Summed to a system total, which sets the tank pressure required to deliver the design inlet condition at the injector:

$$p_{\text{tank}} = p_{\text{injector}} + \sum \Delta p$$

A verification tool for the poppet valves

Alongside the feed-system sizing, I built a second automated calculator for the teammates designing our poppet pneumatic valves, so they could check their own work as they went. It runs the three checks that decide whether a valve body holds pressure and actuates as intended: hoop stress in the body, a force balance across the poppet, and thread shear at the joints. Handing them a tool rather than a worksheet meant a valve geometry could be verified in seconds, and it gave the team one shared reference so everyone was checking against the same math.

Why it mattered

The tool turned line sizing into a trade study rather than a guess. It is also what made the diagnosis on [Project Clementine](#) immediate: when the measured GN2 pressure drop came in high, the model said what the drop *should* have been for that geometry, which pointed straight at undersized tubing instead of leaving us to hunt.

What I would do differently

I would build compressible flow handling in from the start. The incompressible relations are correct for the liquid propellant runs, but the high-pressure GN2 side is better described with compressibility and choked flow limits accounted for,

$$\frac{p^*}{p_0} = \left(\frac{2}{\gamma + 1} \right)^{\frac{\gamma}{\gamma - 1}} \approx 0.528 \quad (\gamma = 1.4)$$

so that the tool covers the pressurant side with the same rigor as the propellant side.

Poppet Valve Verification — Hoop Stress, Force Balance, Thread Shear									
EDIT BLUE VALUES ONLY									
INPUTS				RESULTS / VERIFICATION					
Parameter	Symbol	Units	Cryogenic / LOX	Non-Cryogenic	Metric	Units	What it checks	Cryogenic / LOX	Non-Cryogenic
HOOP STRESS CHECK									
Internal pressure for hoop check	P_int	psi	1000.000	600.000	Hoop stress	psi	$\sigma_H = P_{int} D / (2t)$	4208.333	1949.045
Upstream pressure	P_up	psi	1000.000	600.000	Hoop FOS	-	$\sigma_{allow} / \sigma_H$	4.277	15.392
Pilot pressure	P_pilot	psi	100.000	300.000	Max pressure from hoop	psi	$2t \sigma_{allow} / D$	4277.228	9235.294
Valve inner diameter	D	in	0.805	1.020	FORCE BALANCE CHECK				
Valve wall thickness	t	in	0.090	0.157	Upper area A1	in^2	$\pi r1^2$	0.106	0.754
Upper force radius	r1	in	0.313	0.490	Pilot annular area A2	in^2	$\pi (r2^2 - r1^2)$	2.098	0.000
Pilot force outer radius	r2	in	0.875	0.490	Upstream force F1	lbf	$P_{up} A_1$	106.500	452.578
Effective pressure diameter for thread of head	d_eff	in	0.134	0.375	Pilot force F2	lbf	$P_{pilot} A_2$	209.849	0.000
Mean thread diameter	d_m	in	0.104	0.370	Force ratio F2/F1	-	F_2 / F_1	1.970	0.000
Thread engagement length	L_e	in	0.720	0.190	Force balance status	-	ratio within target $\pm$ tolerance	PASS	REVIEW
Threads per inch	TPI	1/in	32.000	16.000	THREAD SHEAR CHECK				
Allowable normal stress	σ_{allow}	psi	18000.000	30000.000	Effective pressure area	in^2	$\pi d_{eff}^2 / 4$	0.014	0.110
Allowable shear stress	τ_{allow}	psi	12000.000	20000.000	Thread load	lbf	$P_{up} A_{eff}$	14.103	66.268
Required minimum FOS	FOS_req	-	2.000	2.000	Thread shear area	in^2	$\pi d_m L_e$	0.371	0.221
Force ratio target F2/F1	ratio_tgt	-	2.000	2.000	Thread shear stress	psi	F_{thread} / A_{shear}	38.017	300.053
Force ratio tolerance	tol	-	0.100	0.100	Thread FOS	-	$\tau_{allow} / \tau_{thread}$	315.652	66.655
					Max pressure from thread	psi	$\tau_{allow} A_{shear} / A_{eff}$	315651.593	39992.889
					Threads engaged	count	$L_e TPI$	23.040	3.040
					Min threads required	count	$F_{thread} TPI / (\tau_{allow} \pi d_m)$	0.073	0.046
					Thread status	-	engaged $\geq$ required and FOS $\geq$ FOS_req	PASS	PASS
OVERALL									
Overall minimum FOS	-		Min(Hoop FOS, thread FOS)				4.277	15.392	
Overall status	-		force balance PASS and thread PASS and hoop FOS $\geq$ req				PASS	REVIEW	

Figure 2. The poppet-valve verification tool I built for the team: hoop stress, force balance, and thread shear.

Supplier Cost-Benefit Analysis & Validation

Materials sourcing for 3D-printed dental prosthetics

ROLE	ORGANIZATION	PERIOD
Engineering Research & Development Intern	Spectrum Killian Dental Lab Alliance (SKDLA)	Jul 2025 - Oct 2025

The problem

The lab bought a consumable staining material for its 3D-printed dental prosthetics from an established supplier, at a price that had never been challenged. An alternative supplier offered a comparable material for less. But a lower sticker price is not a decision: switching a production material puts product quality at risk, and for a patient-facing medical product quality cannot be allowed to regress. My task was to determine, with evidence, whether the alternative was genuinely equivalent and genuinely worth adopting.

Constraints

- **Quality could not regress.** The end product is worn by patients. A cost saving that compromised quality was not a saving.
- **A live production environment.** Any test had to run without disrupting the technicians' ongoing work.
- **A decision that had to survive scrutiny.** The recommendation would drive real spend, so it had to be defensible to the technicians who use the material and the management who pay for it.
- **Proprietary boundaries.** Supplier terms and internal costs were confidential, so the work had to be documented in a way that captured the method without exposing specifics.

Approach

I treated it as a trade study, not a price comparison. I built the cost side from total cost of use rather than unit price alone, so that anything the cheaper material might cost elsewhere, in waste, rework, or slower processing, would show up rather than hide. On the performance side, I defined what "equivalent quality" had to mean in measurable terms before running anything, so the result could not be argued with after the fact.

What I delivered

- A **side-by-side validation test** that I organized and led, putting the current material and the alternative in front of both the technicians who apply it and the management who judge the result,

under the same conditions. Structuring it as a direct comparison made the outcome evidence rather than opinion.

- A **comparative report** that laid cost and quality out together and made a clear recommendation. This is what the sourcing decision was actually made on.
 - **Standard operating procedures and test protocols**, so the change could be adopted consistently across multiple technicians instead of living in one person's head.
-

Result

The alternative met the quality bar in the side-by-side, and the analysis put a number on the opportunity: more than a 30% material cost reduction, into six figures of projected annual savings. The lab ultimately chose not to switch suppliers, but that call was made on evidence rather than assumption, which was the point of the work. The standard operating procedures I wrote outlasted the question, giving the team a repeatable way to evaluate a material change across technicians.

What I took from it

The habit I built here is the one I now use on propulsion hardware: decide what "good enough" has to mean, in measurable terms, before you run the test, and build the cost case from total cost of use rather than the number on the invoice. The cheapest or highest-performing option on paper is frequently not the right selection once the full picture is in. It is the same discipline behind how I approach component selection on a fluid system.

Bulkhead Drill Jig

Overnight fixture design under a test deadline

ROLE	ORGANIZATION	PERIOD
Designer and machinist	Highlander Space Program, UC Riverside	2025



Figure 1. The drill jig (blue) clamped on the pressure vessel.

The problem

The night before we were due to travel out and hot fire the rocket, the pressure vessel was not prepped. The bulkheads could not be installed because the holes at the ends had not been drilled. Without those holes there was no vessel, and without the vessel there was no test.

Constraints

- **One night.** The team was leaving in the morning.
- **No second chance on the part.** This was the flight pressure vessel. A misplaced hole would scrap a component we could not replace before the test.
- **Repeatability across both ends,** and hole positions accurate enough that the bulkhead fastener pattern would actually line up.

Design decisions

The instinct under time pressure is to mark out the holes by hand and drill. I did not, because hand layout puts the positional accuracy of a flight part in a scribe line and a steady hand, and it has to be repeated independently at both ends. Each layout is a fresh opportunity to be wrong.

A jig moves the accuracy problem off the part and into a fixture. The hole pattern is established once, in the jig, and then transferred identically to every part and every end. The bolt circle positions follow directly from the pattern geometry,

$$\theta_i = \frac{360^\circ}{n} i, \quad x_i = R \cos \theta_i, \quad y_i = R \sin \theta_i$$

where n is the number of fasteners and R the bolt circle radius. Locating off the jig also constrains the drill against wandering at entry, which on a curved surface is the most likely way to lose position.

Spending part of a very short night building a fixture rather than drilling immediately was the decision that mattered. It cost time up front and bought back accuracy and a second identical end.

Outcome

The vessel was drilled, the bulkheads fit on both ends, and the pressure vessel was ready by morning. We travelled on schedule and ran the hot fire.

What I would do differently

The jig should have existed before the night it was needed. The requirement was knowable weeks earlier. What this taught me is to look ahead at which operations on the critical path have no tooling yet, because those are the ones that turn into overnight problems.



Figure 2. The finished drill jig.



Figure 3. Bulkhead holes located on the pressure vessel.

Composite Overwrapped Pressure Vessel Study

Carbon fiber overwrap testing for higher tank pressure

ROLE	ORGANIZATION	PERIOD
General Member	Highlander Space Program, UC Riverside	2024 - 2025



Figure 1. Assembling the COPV test vessel.

The problem

Project Poseidon fed its engine from pressurized propellant tanks. Higher tank pressure raises chamber pressure, and higher chamber pressure raises thrust. What limits tank pressure is the tank wall itself: push a metal vessel too far and the wall stress reaches the material's yield point. You can allow more pressure by making the wall thicker, but that adds weight the vehicle then has to carry all the way up.

The question our small group set out to answer was whether a composite overwrapped pressure vessel, a COPV, could hold higher pressure without that weight penalty. I joined the effort as a general member.

Constraints

- **General-member role.** I was not leading this. I owned the CAD, and I shared the build and test with two teammates.
- **Student fabrication.** The test vessels, bulkheads, and tooling all had to be things we could machine and assemble ourselves.

- **A process we were still learning.** None of us were experienced with carbon fiber layup, so getting the process right was part of what we were testing rather than something we could assume.
- **Pressurized test article.** Anything that holds pressure is a hazard, so the parts had to be built and handled with that in mind.

Design decisions

Why a composite overwrap instead of a thicker metal wall

A thin-walled pressure vessel carries a hoop stress that grows with pressure and radius and falls with wall thickness,

$$\sigma_h = \frac{p r}{t}$$

To allow more pressure in an all-metal vessel you raise t , and that adds mass around the entire vessel. A COPV takes a different route. A thin metal liner holds the gas and provides the seal, while a carbon fiber overwrap carries most of the pressure load in tension. Carbon fiber has a very high strength for its weight, so the overwrap buys pressure capacity for far less mass than the equivalent metal would. That trade is the reason COPVs are used on flight vehicles in the first place.

Capping the vessel with bulkheads and a drill jig

To turn a plain section of tube into a vessel we could actually pressurize, both ends had to be closed. I designed two bulkheads in SolidWorks to cap the ends of a smaller test vessel, and a drill jig to locate the fastener holes.

The jig was the part that mattered most. The holes in the bulkhead have to line up with the holes in the vessel, and drilling them by hand without a guide would not hit the same pattern twice. Locating them with tooling instead of by eye is the same reasoning I used again later on the [Bulkhead Drill Jig](#) for Poseidon.

Building the overwrap by hand

With the vessel assembled, we wrapped carbon fiber around one of the two test vessels and left the other bare as a reference. The goal was to learn the process directly: how the fiber is tensioned and oriented, how the layup goes down, and how the overwrapped vessel behaves under pressure compared to the bare one.

What I built

Two bulkheads and a drill jig, all designed in SolidWorks, that turned a small tube into a sealed, testable pressure vessel. Working with two teammates, I helped assemble the vessels and overwrap one of them in carbon fiber to make the COPV test article.

Validation

We pressurized both articles and compared the bare vessel against the carbon fiber overwrapped one, so we could see directly how much pressure capacity the overwrap added rather than estimating it. That comparison, together with what it took to build and handle a COPV as a student team, is what the program used to weigh whether a COPV belonged on Poseidon.

What I would do differently

I would instrument the pressure test more heavily. We compared the two articles, but with more measurement across the wall we would have understood how the overwrap carried load, not only the point it reached, and we would have had data to design the next article against.

The more honest lesson is about what the project was really for. We did not end up putting a COPV on Poseidon. What we kept was the carbon fiber capability: we used what we learned about composite layup to build lightweight housings for the fluid components exposed underneath the pressure vessel, on the line running to the combustion chamber. The COPV was the goal we set, but the composites skill was the result we actually carried forward, and it went on to fly on [Project Poseidon](#).



Figure 2. The two bulkheads I designed in SolidWorks to cap the test vessel.



Figure 3. Wrapping carbon fiber around the test vessel.




Figure 4. The assembled overwrapped test article.

Formula SAE Suspension Upright

Three design iterations for manufacturability

ROLE	ORGANIZATION	PERIOD
Suspension Engineer	Formula SAE, UC Riverside	2024 - 2025




Figure 1. The final upright design, the third iteration.

The problem

The upright carries wheel loads into the suspension and locates the hub, so it needs tight tolerances on its critical features. The first design achieved that geometry but was expensive to make: complex features and awkward tool access meant more machining operations than the part warranted.

A part that cannot be produced accurately and repeatably with the equipment actually available is not a finished design, regardless of how it performs in CAD.

Design decisions

Across three iterations I worked the geometry against the manufacturing process rather than only against the load case. The levers that mattered:

- **Reducing the number of setups.** Every time a part is unclamped and re-fixtured, cost goes up and a new tolerance stack is introduced between features cut in different orientations. Consolidating features into fewer orientations improves both cost and achievable accuracy.

- **Improving tool access.** Features that force long, thin tooling or unusual approach angles machine slowly and chatter. Opening up access let standard tooling reach the features.
- **Simplifying geometry that was not carrying load.** Complexity that existed for its own sake was removed; complexity that carried wheel loads stayed.

The net result was roughly a 25% improvement in manufacturability by reduced machining complexity and better feature accessibility, while preserving the precision required at the hub and mounting interfaces.

What I would do differently

I would involve whoever is cutting the part in the first iteration rather than the second. Most of what I found in cycles two and three, a machinist would have flagged in five minutes looking at cycle one. Design for manufacture works far better as a conversation than as a solo review.




Figure 2. First iteration in SolidWorks.



Figure 3. Second iteration.